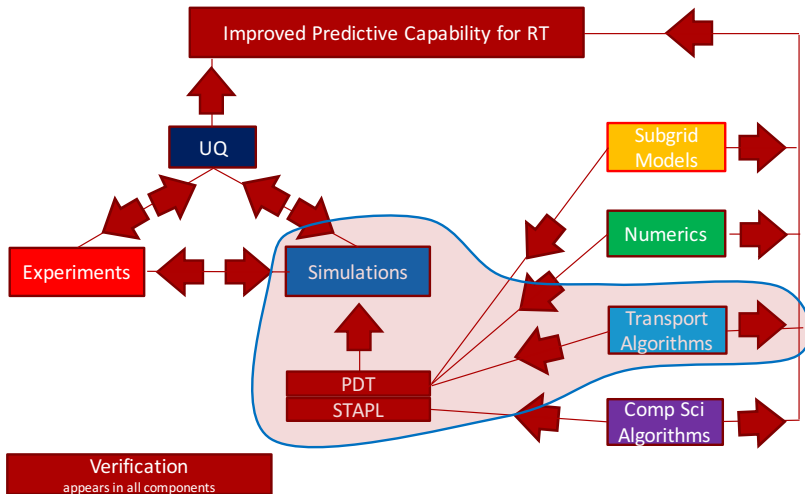


Load Balancing Unstructured Meshes for Massively Parallel Transport Sweeps

Tarek Ghaddar
Dr. Jean Ragusa

Texas A&M University



Project Components and Integration



- 1 Introduction
- 2 Load Balance Algorithm
- 3 Load Balancing Results
- 4 Conclusions


- When running any massively parallel code, load balancing is a priority in order to achieve the best possible parallel efficiency.
- A load balanced problem has an equal number of degrees of freedom per processor.
- Load balancing a logically Cartesian mesh is “not difficult”, as the user specifies the number of cells being used.
- In an unstructured mesh, the user cannot always specify the number of cells they want per processor, and obtaining a load balanced problem is more difficult.
- The goal is to implement a load balancing algorithm for unstructured meshes in PDT.

- Unstructured meshes in PDT are generated in 2D using the Triangle Mesh Generator.
- These can be extruded to create 3D meshes.



- The user inputs coordinates for cut lines in the X and Y directions.
- The cut lines will determine the number of “subsets” the problem is partitioned into.
- Optimizing the location of these cut lines is the basis of the load balancing algorithm.
- A “subset” is an orthogonal unit that is formed by intersecting cut lines.

The Subset



- **Goal:** Obtain an equal number of cells per processor, which for our purposes means an equal number of cells per subset.
- Achieved by optimizing the location of X_i and Y_j , the location of the cut lines.
- **Define:**
 - N_{ij} : The number of cells in subset i, j
 - $f = \frac{\max_{ij}(N_{ij})}{\frac{N_{tot}}{I \cdot J}}$
 - $f_I = \max_i [\sum_j N_{ij}] / \frac{N_{tot}}{I}$
 - $f_J = \max_j [\sum_i N_{ij}] / \frac{N_{tot}}{J}$


```
//Check if all subsets meet the tolerance
while (f > tol_subset)
{
    if (f_l > tol_column)
    {
        Redistribute(X);
    }
    if (f_J > tol_row)
    {
        Redistribute(Y);
    }


    Remesh;
}
```

```
while (f_l > tol_column)
{
    if (f_l > tol_column)
    {
        Redistribute(X);
    }


    Remesh;
}
```

```
while (f > tol)
    if (f_J > tol_row)
    {
        Redistribute(Y);
    }
    Remesh;
}
```

Redistribution Function




Redistribution Function



- Three test cases were used to study the behavior of the load balancing algorithm.
- For each test case, 162 inputs were constructed by varying:
 - The number of subsets
 - The spatial resolution of the mesh (maximum triangle area).

Test Case 1



Test Case 1 - Original Load Balancing Method

1: The best metric behavior of the first test case after **10 load balancing iterations**.

Area	N=4	N=9	N=16	N=25	N=36	N=49	N=64	N=81	N=100
Coarse	1.949	1.598	3.368	2.098	2.278	2.678	2.535	2.805	3.053
1.8	1.458	1.940	2.449	2.590	2.980	3.441	2.957	4.657	3.434
1.6	1.423	1.949	2.427	2.411	3.004	3.053	3.579	4.107	4.105
1.4	1.316	1.871	2.654	3.130	2.451	3.030	3.473	4.040	3.898
1.2	1.298	1.765	2.462	2.656	2.592	3.178	3.144	4.282	4.683
1	1.348	1.638	2.260	2.327	2.347	3.013	3.357	3.841	4.245
0.8	1.257	1.513	2.017	2.792	2.018	2.617	2.884	3.423	3.629
0.6	1.142	1.452	1.788	2.408	2.332	2.092	2.669	2.874	3.629
0.4	1.095	1.353	1.449	1.872	2.397	1.836	2.153	2.351	2.262
0.2	1.046	1.136	1.336	1.545	1.648	2.049	1.678	1.790	1.714
0.1	1.020	1.043	1.109	1.170	1.287	1.357	1.297	1.409	1.221
0.08	1.011	1.029	1.094	1.190	1.209	1.290	1.268	1.318	1.381
0.06	1.005	1.031	1.037	1.105	1.087	1.189	1.177	1.283	1.068
0.05	1.021	1.022	1.058	1.092	1.079	1.115	1.157	1.218	1.176
0.04	1.004	1.013	1.002	1.061	1.074	1.073	1.158	1.171	1.171
0.03	1.003	1.016	1.021	1.050	1.065	1.048	1.928	1.132	1.041
0.02	1.004	1.008	1.010	1.034	1.024	1.028	1.574	1.075	1.094
0.01	1.003	1.010	1.008	1.009	1.039	1.018	1.276	1.043	1.022

Test Case 1 - Load Balancing By Dimension


2: The best metric behavior of the first test case after **10** load balancing by dimension iterations.

Area	N=4	N=9	N=16	N=25	N=36	N=49	N=64	N=81	N=100
Coarse	1.854	1.205	1.475	1.398	1.260	1.367	1.428	1.625	1.639
1.800	1.079	1.392	1.523	1.722	2.036	2.315	2.421	3.962	3.587
1.600	1.044	1.704	1.536	1.566	1.791	2.126	2.492	3.417	3.807
1.400	1.073	1.152	1.509	1.756	1.811	2.183	2.475	2.911	3.974
1.200	1.037	1.125	1.515	1.718	1.972	2.611	2.628	3.816	3.205
1.000	1.047	1.209	1.180	1.763	1.494	1.928	2.564	3.744	3.953
0.800	1.077	1.193	1.246	1.343	1.417	2.949	2.411	3.560	3.629
0.600	1.066	1.112	1.388	1.455	1.521	1.837	1.769	2.586	3.381
0.400	1.095	1.032	1.145	1.221	1.479	1.355	1.554	2.038	1.838
0.200	1.046	1.049	1.075	1.149	1.173	1.174	1.210	1.251	1.571
0.100	1.020	1.043	1.023	1.071	1.103	1.138	1.129	1.091	1.205
0.080	1.011	1.029	1.094	1.072	1.076	1.097	1.068	1.071	1.213
0.060	1.005	1.031	1.037	1.016	1.087	1.096	1.093	1.095	1.068
0.050	1.021	1.022	1.058	1.092	1.079	1.062	1.090	1.075	1.088
0.040	1.004	1.013	1.002	1.061	1.074	1.073	1.067	1.090	1.100
0.030	1.003	1.016	1.021	1.050	1.065	1.048	1.038	1.061	1.041
0.020	1.004	1.008	1.010	1.034	1.024	1.028	1.058	1.075	1.094
0.010	1.003	1.010	1.008	1.009	1.039	1.018	1.054	1.043	1.022

3: The improvement with the load balancing by dimension method.

Area	N=4	N=9	N=16	N=25	N=36	N=49	N=64	N=81	N=100
Coarse	0.049	0.246	0.562	0.334	0.447	0.489	0.437	0.421	0.463
1.800	0.260	0.282	0.378	0.335	0.317	0.327	0.181	0.149	-0.045
1.600	0.266	0.126	0.367	0.350	0.404	0.304	0.304	0.168	0.072
1.400	0.185	0.384	0.431	0.439	0.261	0.279	0.287	0.279	-0.020
1.200	0.201	0.363	0.384	0.353	0.239	0.178	0.164	0.109	0.316
1.000	0.223	0.262	0.478	0.242	0.364	0.360	0.236	0.025	0.069
0.800	0.143	0.211	0.382	0.519	0.298	-0.127	0.164	-0.040	0.000
0.600	0.067	0.234	0.224	0.396	0.347	0.122	0.337	0.100	0.068
0.400	0.000	0.237	0.210	0.347	0.383	0.262	0.278	0.133	0.188
0.200	0.000	0.076	0.196	0.256	0.288	0.427	0.279	0.301	0.083
0.100	0.000	0.000	0.078	0.085	0.143	0.161	0.130	0.226	0.013
0.080	0.000	0.000	0.000	0.099	0.110	0.150	0.158	0.188	0.122
0.060	0.000	0.000	0.000	0.080	0.000	0.078	0.071	0.147	0.000
0.050	0.000	0.000	0.000	0.000	0.000	0.048	0.058	0.117	0.075
0.040	0.000	0.000	0.000	0.000	0.000	0.000	0.079	0.069	0.061
0.030	0.000	0.000	0.000	0.000	0.000	0.000	0.462	0.062	0.000
0.020	0.000	0.000	0.000	0.000	0.000	0.000	0.328	0.000	0.000
0.010	0.000	0.000	0.000	0.000	0.000	0.000	0.174	0.000	0.000

Test Case 2



4: The best metric behavior of the second test case after **10 load balancing iterations**.

Area	N=4	N=9	N=16	N=25	N=36	N=49	N=64	N=81	N=100
Coarse	1.854	1.361	1.765	1.479	1.742	1.595	1.792	1.820	1.923
1.800	1.176	1.336	1.882	2.375	2.269	2.359	2.544	3.841	4.874
1.600	1.109	1.482	1.783	1.701	1.990	2.421	2.848	3.345	2.989
1.400	1.133	1.366	1.854	1.746	1.882	2.600	2.833	3.617	2.692
1.200	1.153	1.506	1.575	1.599	2.162	2.270	2.562	3.355	3.771
1.000	1.132	1.418	1.729	1.694	1.581	2.452	2.491	3.231	3.902
0.800	1.139	1.355	1.437	1.610	1.940	2.167	2.152	2.250	2.936
0.600	1.053	1.360	1.604	1.705	1.687	1.960	1.902	2.458	2.500
0.400	1.095	1.176	1.401	1.534	1.771	1.797	1.841	1.792	2.262
0.200	1.043	1.140	1.183	1.364	1.561	1.741	1.587	1.495	1.714
0.100	1.028	1.042	1.114	1.193	1.284	1.335	1.268	1.227	1.283
0.080	1.013	1.037	1.091	1.190	1.210	1.293	1.197	1.236	1.178
0.060	1.007	1.033	1.037	1.105	1.087	1.205	1.183	1.281	1.101
0.050	1.021	1.026	1.050	1.088	1.061	1.115	1.182	1.215	1.176
0.040	1.005	1.019	1.001	1.061	1.075	1.073	1.098	1.173	1.171
0.030	1.005	1.013	1.021	1.045	1.060	1.050	1.265	1.101	1.041
0.020	1.006	1.017	1.008	1.034	1.022	1.024	1.186	1.076	1.097
0.010	1.003	1.010	1.008	1.009	1.039	1.018	1.276	1.043	1.022


Test Case 2 - Load Balancing By Dimension

5: The best metric behavior of the second test case after **10 load balancing by dimension iterations**.


Area	N=4	N=9	N=16	N=25	N=36	N=49	N=64	N=81	N=100
Coarse	1.854	1.134	1.672	1.324	1.653	1.518	1.542	1.401	1.520
1.800	1.099	1.135	1.438	1.868	1.753	1.840	2.740	3.262	1.690
1.600	1.086	1.118	1.598	1.558	1.802	2.350	4.700	3.776	3.441
1.400	1.052	1.101	1.432	1.768	1.966	2.000	2.688	2.827	3.169
1.200	1.042	1.183	1.419	1.542	1.620	2.192	3.939	3.237	3.111
1.000	1.057	1.207	1.315	1.331	1.453	1.698	2.436	2.455	3.524
0.800	1.051	1.097	1.204	1.570	1.510	1.753	1.738	2.232	2.576
0.600	1.053	1.091	1.070	1.215	1.424	1.765	1.613	1.678	2.456
0.400	1.095	1.087	1.120	1.212	1.226	1.346	1.215	2.078	2.128
0.200	1.043	1.038	1.074	1.112	1.238	1.082	1.311	1.564	1.522
0.100	1.028	1.042	1.097	1.042	1.098	1.124	1.129	1.095	1.204
0.080	1.013	1.037	1.091	1.085	1.090	1.128	1.092	1.110	1.178
0.060	1.007	1.033	1.037	1.014	1.087	1.082	1.048	1.092	1.034
0.050	1.021	1.026	1.050	1.088	1.061	1.052	1.083	1.079	1.075
0.040	1.005	1.019	1.001	1.061	1.075	1.073	1.081	1.093	1.145
0.030	1.005	1.013	1.021	1.045	1.060	1.050	1.061	1.076	1.041
0.020	1.006	1.017	1.008	1.034	1.022	1.024	1.095	1.076	1.097
0.010	1.003	1.010	1.008	1.009	1.039	1.018	1.092	1.043	1.022

6: The improvement with the load balancing by dimension method.


Area	N=4	N=9	N=16	N=25	N=36	N=49	N=64	N=81	N=100
Coarse	0.000	0.167	0.053	0.105	0.051	0.049	0.139	0.230	0.210
1.800	0.066	0.151	0.236	0.214	0.228	0.220	-0.077	0.151	0.653
1.600	0.021	0.246	0.104	0.084	0.094	0.030	-0.650	-0.129	-0.151
1.400	0.072	0.193	0.227	-0.012	-0.045	0.231	0.051	0.218	-0.177
1.200	0.096	0.215	0.099	0.036	0.251	0.035	-0.537	0.035	0.175
1.000	0.066	0.149	0.239	0.214	0.081	0.308	0.022	0.240	0.097
0.800	0.078	0.191	0.162	0.025	0.221	0.191	0.192	0.008	0.123
0.600	0.000	0.198	0.333	0.288	0.156	0.099	0.152	0.318	0.018
0.400	0.000	0.075	0.201	0.210	0.308	0.251	0.340	-0.159	0.060
0.200	0.000	0.089	0.092	0.185	0.207	0.379	0.174	-0.046	0.112
0.100	0.000	0.000	0.015	0.126	0.145	0.158	0.109	0.108	0.062
0.080	0.000	0.000	0.000	0.089	0.099	0.128	0.088	0.102	0.000
0.060	0.000	0.000	0.000	0.082	0.000	0.102	0.115	0.148	0.061
0.050	0.000	0.000	0.000	0.000	0.000	0.057	0.084	0.112	0.086
0.040	0.000	0.000	0.000	0.000	0.000	0.000	0.016	0.069	0.022
0.030	0.000	0.000	0.000	0.000	0.000	0.000	0.162	0.022	0.000
0.020	0.000	0.000	0.000	0.000	0.000	0.000	0.077	0.000	0.000
0.010	0.000	0.000	0.000	0.000	0.000	0.000	0.144	0.000	0.000



IM1 - Original Load Balancing, $f = 2.63$



IM1 - Load Balancing By Dimension, $f = 1.76$



- The effectiveness of the load balancing algorithm depends on the spatial distribution of fine geometric features, the maximum triangle area used, and the number of subsets the domain is decomposed into.
- Good improvement is seen for Test Cases 1 and 2, and especially the IM1 problem.m
- More tinkering with the load balancing by dimension algorithm will be done to study its behavior and potential improvements.

- Moving away from the Triangle Mesh Generator
 - Lack of support
 - Unable to enforce mesh quality consistently while load balancing problems
- Current path forward involves a collaborative effort with Richard Vega (TAMU/Sandia) using a combination of Cubit and OpenFoam.
 - Splitting the mesh into subsets rather than the problem geometry

- Two more paths for improving the load balancing algorithm have been outlined.
 - Adaptively splitting the subsets that have large cell counts into smaller subsets, and redistributing subsets amongst processors.
 - Taking advantage of nested parallelism to assign more parallel processes at subsets that require more work to be done.
- Studying the behavior of the communication penalty while load balancing (brought up by Derek Gaston at M&C 2017).

Initial Setup

P:2 4	P:5 16	P:8 4
P:1 16	P:4 4	P:7 16
P:0 4	P:3 4	P:6 4

P: Processor_ID
Number of Cells

$$f = 2$$

Adaptively Refine Subsets

P:2 4		P:5 4	P:5 4	P:8 4	
		P:5 4	P:5 4		
P:1 4	P:1 4	P:4 4		P:7 4	P:7 4
P:1 4	P:1 4			P:7 4	P:7 4
P:0 4		P:3 4		P:6 4	

Subset Redistribution

$$f = 1$$

P:2 4		P:5 4	P:5 4	P:8 4	
		P:2 4	P:4 4		
P:1 4	P:1 4	P:4 4		P:8 4	P:7 4
P:0 4	P:3 4			P:6 4	P:7 4
P:0 4		P:3 4		P:6 4	

Nested Parallelism

P:2 4	P:5 16	P:8 4
P:1 16	P:4 4	P:7 16
P:0 4	P:3 4	P:6 4

P: Processor_ID
Number of Cells

$$f = 2$$

A special thank you to the following individuals for their help and support:

- Drs. Ragusa, Morel, Adams, and Popov
- Michael Adams, Daryl Hawkins, and Dr. Timmie Smith
- Dr. Andrew Till
- The CERT team and fellow grad students
- PSAAP-II